

anime-face-generation

March 30, 2024

0.1 Import Modules

```
[ ]: import os
import numpy as np
import matplotlib.pyplot as plt
import warnings
from tqdm.notebook import tqdm

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import load_img, array_to_img
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras import layers
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import BinaryCrossentropy

warnings.filterwarnings('ignore')
```

0.2 Load the files

```
[ ]: BASE_DIR = '/kaggle/input/anime-faces/data/'
```

```
[ ]: # load complete image paths to the list
image_paths = []
for image_name in os.listdir(BASE_DIR):
    image_path = os.path.join(BASE_DIR, image_name)
    image_paths.append(image_path)
```

```
[ ]: image_paths[:5]
```

```
[ ]: ['/kaggle/input/anime-faces/data/21130.png',
      '/kaggle/input/anime-faces/data/9273.png',
      '/kaggle/input/anime-faces/data/18966.png',
      '/kaggle/input/anime-faces/data/14127.png',
      '/kaggle/input/anime-faces/data/18054.png']
```

```
[ ]: # remove unnecessary file
image_paths.remove('/kaggle/input/anime-faces/data/data')
```

```
[ ]: len(image_paths)
```

```
[ ]: 21551
```

0.3 Visualize the Image Dataset

```
[ ]: # to display grid of images (7x7)
plt.figure(figsize=(20, 20))
temp_images = image_paths[:49]
index = 1

for image_path in temp_images:
    plt.subplot(7, 7, index)
    # load the image
    img = load_img(image_path)
    # convert to numpy array
    img = np.array(img)
    # show the image
    plt.imshow(img)
    plt.axis('off')
    # increment the index for next image
    index += 1
```



0.4 Preprocess Images

```
[ ]: # load the image and convert to numpy array
train_images = [np.array(load_img(path)) for path in tqdm(image_paths)]
train_images = np.array(train_images)
```

```
0%|          | 0/21551 [00:00<?, ?it/s]
```

```
[ ]: train_images[0].shape
```

```
[ ]: (64, 64, 3)
```

```
[ ]: # reshape the array
train_images = train_images.reshape(train_images.shape[0], 64, 64, 3).
↳astype('float32')
```

```
[ ]: # normalize the images
train_images = (train_images - 127.5) / 127.5
```

```
[ ]: train_images[0]
```

```
[ ]: array([[[-0.7254902 , -0.90588236, -0.7176471 ],
             [-0.70980394, -0.88235295, -0.6862745 ],
             [-0.7019608 , -1.          , -0.7176471 ],
             ...,
             [-0.09019608,  0.01176471,  0.3882353 ],
             [-0.19215687,  0.05882353,  0.4509804 ],
             [-0.34117648, -0.13725491,  0.29411766]],

            [[[-0.7254902 , -0.90588236, -0.7176471 ],
             [-0.73333335, -0.8666667 , -0.7019608 ],
             [-0.52156866, -0.8666667 , -0.5686275 ],
             ...,
             [-0.29411766,  0.01176471,  0.3019608 ],
             [-0.08235294,  0.15294118,  0.5294118 ],
             [-0.23921569,  0.09019608,  0.45882353]],

            [[[-0.7019608 , -0.8901961 , -0.7019608 ],
             [-0.7176471 , -0.8745098 , -0.69411767],
             [-0.5686275 , -0.8509804 , -0.58431375],
             ...,
             [-0.22352941,  0.1764706 ,  0.48235294],
             [-0.13725491,  0.1764706 ,  0.4745098 ],
             [-0.06666667,  0.24705882,  0.5058824 ]],

            ...,

            [[[-0.30980393, -0.6862745 , -0.29411766],
             [-0.34901962, -0.6862745 , -0.37254903],
             [-0.30980393, -0.5372549 , -0.35686275],
             ...,
             [-0.79607844, -0.9843137 , -0.77254903],
             [-0.827451   , -0.8745098 , -0.85882354],
             [-0.7647059 , -0.92156863, -0.8039216 ]],

            [[[-0.3254902 , -0.6392157 , -0.29411766],
             [-0.44313726, -0.7254902 , -0.43529412],
             [-0.54509807, -0.7882353 , -0.54509807],
             ...,
```

```

[-0.8352941 , -0.9764706 , -0.70980394],
[-0.7254902 , -0.8745098 , -0.75686276],
[-0.7411765 , -0.9137255 , -0.69411767]],

[[-0.5764706 , -0.827451 , -0.54509807],
 [-0.56078434, -0.8666667 , -0.5294118 ],
 [-0.58431375, -0.85882354, -0.5764706 ],
 ...,
 [-0.7411765 , -0.8666667 , -0.6156863 ],
 [-0.7254902 , -0.9372549 , -0.70980394],
 [-0.7647059 , -0.92941177, -0.7176471 ]]], dtype=float32)

```

0.5 Create Generator & Discriminator Models

```

[ ]: # latent dimension for random noise
LATENT_DIM = 100
# weight initializer
WEIGHT_INIT = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)
# no. of channels of the image
CHANNELS = 3 # for gray scale, keep it as 1

```

0.5.1 Generator Model

Generator Model will create new images similar to training data from random noise

```

[ ]: model = Sequential(name='generator')

# 1d random noise
model.add(layers.Dense(8 * 8 * 512, input_dim=LATENT_DIM))
# model.add(layers.BatchNormalization())
model.add(layers.ReLU())

# convert 1d to 3d
model.add(layers.Reshape((8, 8, 512)))

# upsample to 16x16
model.add(layers.Conv2DTranspose(256, (4, 4), strides=(2, 2), padding='same',
    ↪kernel_initializer=WEIGHT_INIT))
# model.add(layers.BatchNormalization())
model.add(layers.ReLU())

# upsample to 32x32
model.add(layers.Conv2DTranspose(128, (4, 4), strides=(2, 2), padding='same',
    ↪kernel_initializer=WEIGHT_INIT))
# model.add(layers.BatchNormalization())
model.add(layers.ReLU())

```

```
# upsample to 64x64
model.add(layers.Conv2DTranspose(64, (4, 4), strides=(2, 2), padding='same',
    ↪kernel_initializer=WEIGHT_INIT))
# model.add(layers.BatchNormalization())
model.add(layers.ReLU())

model.add(layers.Conv2D(CHANNELS, (4, 4), padding='same', activation='tanh'))

generator = model
generator.summary()
```

```
2023-02-08 02:24:22.261035: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2023-02-08 02:24:22.352832: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2023-02-08 02:24:22.353639: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2023-02-08 02:24:22.355500: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
2023-02-08 02:24:22.355865: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2023-02-08 02:24:22.356549: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
2023-02-08 02:24:22.357298: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA
```

Model: "generator"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32768)	3309568

re_lu (ReLU)	(None, 32768)	0
reshape (Reshape)	(None, 8, 8, 512)	0
conv2d_transpose (Conv2DTran	(None, 16, 16, 256)	2097408
re_lu_1 (ReLU)	(None, 16, 16, 256)	0
conv2d_transpose_1 (Conv2DTr	(None, 32, 32, 128)	524416
re_lu_2 (ReLU)	(None, 32, 32, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 64, 64, 64)	131136
re_lu_3 (ReLU)	(None, 64, 64, 64)	0
conv2d (Conv2D)	(None, 64, 64, 3)	3075

=====
Total params: 6,065,603

Trainable params: 6,065,603

Non-trainable params: 0

node zero

2023-02-08 02:24:24.522023: I

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2023-02-08 02:24:24.523055: I

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2023-02-08 02:24:24.523834: I

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:937] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2023-02-08 02:24:24.524488: I

tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 15401 MB memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0

0.5.2 Discriminator Model

Discriminator model will classify the image from the generator to check whether it real (or) fake images.

```
[ ]: model = Sequential(name='discriminator')
input_shape = (64, 64, 3)
```

```

alpha = 0.2

# create conv layers
model.add(layers.Conv2D(64, (4, 4), strides=(2, 2), padding='same',
    ↪input_shape=input_shape))
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU(alpha=alpha))

model.add(layers.Conv2D(128, (4, 4), strides=(2, 2), padding='same',
    ↪input_shape=input_shape))
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU(alpha=alpha))

model.add(layers.Conv2D(128, (4, 4), strides=(2, 2), padding='same',
    ↪input_shape=input_shape))
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU(alpha=alpha))

model.add(layers.Flatten())
model.add(layers.Dropout(0.3))

# output class
model.add(layers.Dense(1, activation='sigmoid'))

discriminator = model
discriminator.summary()

```

Model: "discriminator"

Layer (type)	Output Shape	Param #

conv2d_1 (Conv2D)	(None, 32, 32, 64)	3136

batch_normalization (BatchNo	(None, 32, 32, 64)	256

leaky_re_lu (LeakyReLU)	(None, 32, 32, 64)	0

conv2d_2 (Conv2D)	(None, 16, 16, 128)	131200

batch_normalization_1 (Batch	(None, 16, 16, 128)	512

leaky_re_lu_1 (LeakyReLU)	(None, 16, 16, 128)	0

conv2d_3 (Conv2D)	(None, 8, 8, 128)	262272

batch_normalization_2 (Batch	(None, 8, 8, 128)	512

leaky_re_lu_2 (LeakyReLU)	(None, 8, 8, 128)	0

flatten (Flatten)	(None, 8192)	0

dropout (Dropout)	(None, 8192)	0

dense_1 (Dense)	(None, 1)	8193
=====		
Total params: 406,081		
Trainable params: 405,441		
Non-trainable params: 640		

0.6 Create DCGAN

```
[ ]: class DCGAN(keras.Model):
    def __init__(self, generator, discriminator, latent_dim):
        super().__init__()
        self.generator = generator
        self.discriminator = discriminator
        self.latent_dim = latent_dim
        self.g_loss_metric = keras.metrics.Mean(name='g_loss')
        self.d_loss_metric = keras.metrics.Mean(name='d_loss')

    @property
    def metrics(self):
        return [self.g_loss_metric, self.d_loss_metric]

    def compile(self, g_optimizer, d_optimizer, loss_fn):
        super(DCGAN, self).compile()
        self.g_optimizer = g_optimizer
        self.d_optimizer = d_optimizer
        self.loss_fn = loss_fn

    def train_step(self, real_images):
        # get batch size from the data
        batch_size = tf.shape(real_images)[0]
        # generate random noise
        random_noise = tf.random.normal(shape=(batch_size, self.latent_dim))

        # train the discriminator with real (1) and fake (0) images
        with tf.GradientTape() as tape:
            # compute loss on real images
            pred_real = self.discriminator(real_images, training=True)
            # generate real image labels
            real_labels = tf.ones((batch_size, 1))
            # label smoothing
```

```

        real_labels += 0.05 * tf.random.uniform(tf.shape(real_labels))
        d_loss_real = self.loss_fn(real_labels, pred_real)

        # compute loss on fake images
        fake_images = self.generator(random_noise)
        pred_fake = self.discriminator(fake_images, training=True)
        # generate fake labels
        fake_labels = tf.zeros((batch_size, 1))
        d_loss_fake = self.loss_fn(fake_labels, pred_fake)

        # total discriminator loss
        d_loss = (d_loss_real + d_loss_fake) / 2

        # compute discriminator gradients
        gradients = tape.gradient(d_loss, self.discriminator.
↪ trainable_variables)
        # update the gradients
        self.d_optimizer.apply_gradients(zip(gradients, self.discriminator.
↪ trainable_variables))

        # train the generator model
        labels = tf.ones((batch_size, 1))
        # generator want discriminator to think that fake images are real
        with tf.GradientTape() as tape:
            # generate fake images from generator
            fake_images = self.generator(random_noise, training=True)
            # classify images as real or fake
            pred_fake = self.discriminator(fake_images, training=True)
            # compute loss
            g_loss = self.loss_fn(labels, pred_fake)

        # compute gradients
        gradients = tape.gradient(g_loss, self.generator.trainable_variables)
        # update the gradients
        self.g_optimizer.apply_gradients(zip(gradients, self.generator.
↪ trainable_variables))

        # update states for both models
        self.d_loss_metric.update_state(d_loss)
        self.g_loss_metric.update_state(g_loss)

        return {'d_loss': self.d_loss_metric.result(), 'g_loss': self.
↪ g_loss_metric.result()}

```

```
[ ]: class DCGANMonitor(keras.callbacks.Callback):
    def __init__(self, num_imgs=25, latent_dim=100):
        self.num_imgs = num_imgs
        self.latent_dim = latent_dim
        # create random noise for generating images
        self.noise = tf.random.normal([25, latent_dim])

    def on_epoch_end(self, epoch, logs=None):
        # generate the image from noise
        g_img = self.model.generator(self.noise)
        # denormalize the image
        g_img = (g_img * 127.5) + 127.5
        g_img.numpy()

        fig = plt.figure(figsize=(8, 8))
        for i in range(self.num_imgs):
            plt.subplot(5, 5, i+1)
            img = array_to_img(g_img[i])
            plt.imshow(img)
            plt.axis('off')
        # plt.savefig('epoch_{:03d}.png'.format(epoch))
        plt.show()

    def on_train_end(self, logs=None):
        self.model.generator.save('generator.h5')
```

```
[ ]: dcgan = DCGAN(generator=generator, discriminator=discriminator,
    ↪latent_dim=LATENT_DIM)
```

```
[ ]: D_LR = 0.0001
    G_LR = 0.0003
    dcgan.compile(g_optimizer=Adam(learning_rate=G_LR, beta_1=0.5),
    ↪d_optimizer=Adam(learning_rate=D_LR, beta_1=0.5),
    ↪loss_fn=BinaryCrossentropy())
```


```
[ ]: N_EPOCHS = 50
    dcgan.fit(train_images, epochs=N_EPOCHS, callbacks=[DCGANMonitor()])
```

```
2023-02-08 02:24:25.158449: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 1059274752
exceeds 10% of free system memory.
2023-02-08 02:24:26.332743: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 1059274752
exceeds 10% of free system memory.
2023-02-08 02:24:27.173692: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)
```

Epoch 1/50


2023-02-08 02:24:29.703314: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369]
Loaded cuDNN version 8005

674/674 [=====] - 36s 41ms/step - d_loss: 0.6092 -
g_loss: 2.7299



Epoch 2/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.5241 -
g_loss: 1.9957



Epoch 3/50
674/674 [=====] - 27s 41ms/step - d_loss: 0.5738 -
g_loss: 1.8336



Epoch 4/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.5997 -
g_loss: 1.7399



Epoch 5/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6196 -
g_loss: 1.6546



Epoch 6/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.5933 -
g_loss: 1.5709



Epoch 7/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6225 -
g_loss: 1.3766



Epoch 8/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6364 -
g_loss: 1.2569



Epoch 9/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6491 -
g_loss: 1.1808



Epoch 10/50
674/674 [=====] - 27s 40ms/step - d_loss: 0.6962 -
g_loss: 1.0153



Epoch 11/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.7165 -
g_loss: 0.9253



Epoch 12/50
674/674 [=====] - 27s 41ms/step - d_loss: 0.7209 -
g_loss: 0.8608



Epoch 13/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.7201 -
g_loss: 0.8332



Epoch 14/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.7134 -
g_loss: 0.8114



Epoch 15/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.7064 -
g_loss: 0.8128



Epoch 16/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.7044 -
g_loss: 0.8055



Epoch 17/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.7050 -
g_loss: 0.7948



Epoch 18/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.7024 -
g_loss: 0.7843



Epoch 19/50
674/674 [=====] - 27s 41ms/step - d_loss: 0.7033 -
g_loss: 0.7805



Epoch 20/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.7027 -
g_loss: 0.7699



Epoch 21/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.6987 -
g_loss: 0.7638



Epoch 22/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6996 -
g_loss: 0.7681



Epoch 23/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6940 -
g_loss: 0.7724



Epoch 24/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6930 -
g_loss: 0.7711



Epoch 25/50
674/674 [=====] - 27s 41ms/step - d_loss: 0.6907 -
g_loss: 0.7773



Epoch 26/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6867 -
g_loss: 0.7815



Epoch 27/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6836 -
g_loss: 0.7940



Epoch 28/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6790 -
g_loss: 0.8047



Epoch 29/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6776 -
g_loss: 0.8097



Epoch 30/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.6717 -
g_loss: 0.8222



Epoch 31/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6688 -
g_loss: 0.8341



Epoch 32/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6618 -
g_loss: 0.8468



Epoch 33/50
674/674 [=====] - 28s 41ms/step - d_loss: 0.6582 -
g_loss: 0.8615



Epoch 34/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6516 -
g_loss: 0.8705



Epoch 35/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6470 -
g_loss: 0.8930



Epoch 36/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6424 -
g_loss: 0.9033



Epoch 37/50
674/674 [=====] - 27s 41ms/step - d_loss: 0.6357 -
g_loss: 0.9156



Epoch 38/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6375 -
g_loss: 0.9255



Epoch 39/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6327 -
g_loss: 0.9445




Epoch 40/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6304 -
g_loss: 0.9496



Epoch 41/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6237 -
g_loss: 0.9649



Epoch 42/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6227 -
g_loss: 0.9827



Epoch 43/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6221 -
g_loss: 0.9904



Epoch 44/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.6144 -
g_loss: 0.9968



Epoch 45/50

674/674 [=====] - 28s 41ms/step - d_loss: 0.6125 -
g_loss: 1.0184



Epoch 46/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6076 -
g_loss: 1.0264



Epoch 47/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.6028 -
g_loss: 1.0532



Epoch 48/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.5999 -
g_loss: 1.0589



Epoch 49/50

674/674 [=====] - 27s 41ms/step - d_loss: 0.5981 -
g_loss: 1.0709



Epoch 50/50

674/674 [=====] - 27s 40ms/step - d_loss: 0.5894 -
g_loss: 1.0865



[]: <keras.callbacks.History at 0x7f2f96148c10>

0.7 Generate New Anime Image

```
[ ]: noise = tf.random.normal([1, 100])
fig = plt.figure(figsize=(3, 3))
# generate the image from noise
g_img = dcgan.generator(noise)
# denormalize the image
g_img = (g_img * 127.5) + 127.5
g_img.numpy()
img = array_to_img(g_img[0])
plt.imshow(img)
```



```
plt.axis('off')  
# plt.savefig('epoch_{:03d}.png'.format(epoch))  
plt.show()
```



[]:

[]: